

# Design Experiences in C++

---

## *Designing With Objects*

---

Mark Radford

**twoNine**  
Computer Services Limited

*mark.radford@twonine.co.uk*  
*www.twonine.co.uk*

© twoNine Computer Services Ltd, April 2003

1

## Contents

- Values & quantities in C++
- Choosing an approach to iteration
- Mixing generic with object oriented programming
- Observations on design

2

## Values & Quantities in C++

- *Whole Values*
  - The problem
  - The *Whole Value* solution and its costs
- Quantities
  - Selected issues and pitfalls

3

### The Problem

- In languages without data abstraction support, built in types are used to represent values, consequently:
  - Compile-time type checking is weak
  - Communication is weak (because the domain vocabulary is absent from the code)
- Unfortunately, C++ programmers have a tendency to follow suit

4

## Whole Values

```
class minutes
{
public:
    explicit minutes(int initial);
    ...
};

class transaction
{
public:
    void timeout(
        const minutes& duration);
    ...
};

void f(transaction* current)
{
    current->timeout(minutes(10));
    ...
}
```

- In C++ classes should be used to implement first class data abstractions
  - The compiler is empowered to do stronger type checking
  - The code communicates application domain vocabulary

5

## The Cost

- There is effort in producing the classes
  - Automating their production (e.g. using templates) helps keep this cost down
- The proliferation of classes must be managed
  - This cost is something of a red herring with the right tools
  - Tools cost money, but the cost is tangible

6

## Automation Using Templates – Pitfall

### The problem

```
template <typename T> class value { ... };  
typedef value<unsigned int> seconds;  
typedef value<unsigned int> minutes;
```

Both types are  
the same in  
the eyes of the  
compiler

### Solution by type tagging

```
template <  
    typename type,  
    typename tag> class value { ... };  
  
struct secs_tag {};  
typedef value<int, secs_tag> seconds;  
  
struct mins_tag {};  
typedef value<int, mins_tag> minutes;
```

The empty tag  
structs ensure  
that **seconds**  
and **minutes**  
are both  
unique types

7

## Quantities

- E.g. minutes, voltages, kilograms
  - Quantities account for much of the value based bandwidth in programmes
- Quantities have their own additional specific characteristics, e.g...
  - They have *units*
  - They can have a value of zero

8

## Quantity Abstraction Pitfall

```
class time_interval
{
public:
    struct from_seconds {};
    struct from_minutes {};

    time_interval(
        int value, const from_sec&);
    time_interval(
        int value, const from_min&);
    ...
};

std::ostream& operator<<(
    std::ostream& os,
    const time_interval& value);
```

Not only does  
initialisation  
require a very  
artificial looking  
interface  
design...

*What exactly, is  
inserted into the  
stream?*

9

## Abstract Quantities as Units

```
class minutes
{
public:
    explicit minutes(int initial);
    minutes();
    ...
};

std::ostream& operator<<(
    std::ostream& os,
    const minutes& value);
```

Initialisation is now  
straightforward –  
either by default or  
from a single value

Operator overloads  
come naturally, and  
their meaning is  
automatically clear

10

## Initialisation & Conversion

```
class minutes
{
public:
    typedef int value_type;

    minutes();
    explicit minutes(value_type initial_value);
    value_type value() const;
    ...
};
```

Default initialisation value of zero is a feature of quantity types

Conversion *must* be explicit to convey meaning

Conversion back to underlying type must be invoked *deliberately*

*Note: the provision of value() is necessary for the range of supported operations to be extensible*

11

## The Costs vs. Benefits Summary

- Compile-time error detection is an asset
  - The cost of run-time debugging is open-ended
  - The benefit of compile time error detection is the greater *predictability* of costs
- Whole Value classes belong to the application domain
  - Subsequent projects can build on the existing application domain library

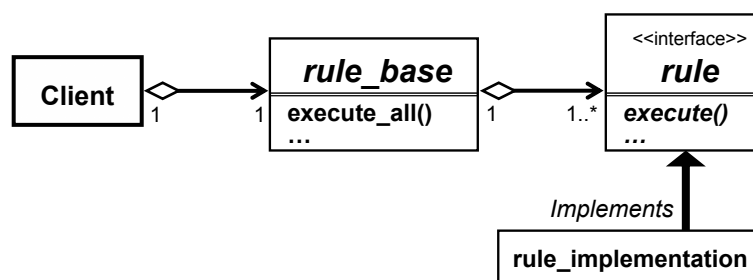
12

## Choosing an Approach to Iteration

- Design of a system of business rules
- Approaches to iteration – two alternatives to the STL approach compared
- A design decision re-evaluated

13

### Example – A System of Business Rules



The **rule\_base** is the repository that manages the storage and execution of the rules.

14

## The Iteration Problem

- `rule_base` must provide some mechanism to allow iteration through the `rules` it stores
- The STL style is not always an appropriate choice
  - It is fragile in cases where updates and iteration might happen concurrently
  - STL style components provide a toolkit, but sometimes more of an application *framework* is needed

15

## Alternatives to STL style iteration

- Batched copying
  - `rule_base` (internally) makes a complete copy of its collection of `rules` and returns it in a single operation
- Enumeration by callback
  - `rule_base` drives the flow of control during iteration
  - A callback function supplied by the client is called for each `rule`

16



## Batched Copying

```
class rule_base
{
public:
    template <
        typename out_it,
        typename copy_fn>
    void copy(
        out_it dest,
        copy_fn make_copy) const
    {
        mutex guard;
        std::transform(
            rules.begin(),
            rules.end(),
            dest,
            make_copy);
    }
    ...
private:
    std::list<rule*> rules;
};
```

- The `copy()` member copies the state of each rule to a sequence starting at `dest`
- Knowledge of the value based data structure used by the sequence is encapsulated in `make_copy()`
- Allowing `rule_base` to drive the control flow allows synchronisation to be made integral to the copying process

17

## Enumeration by Callback

```
class rule_base
{
    typedef std::list<rule*> container;
    typedef container::const_iterator
        const_iterator;
public:
    void for_each(
        void(*callback)(const rule&))
    {
        mutex guard;
        for (const_iterator it=rules.begin();
            it != rules.end();
            ++it)
            callback(*it);
    }
    ...
private:
    container rules;
};
```

- The `for_each()` member calls `callback()` for each stored rule
- Putting `rule_base` in charge of the control flow allows synchronisation to be made integral to the iteration

18

## Interface Class as Callback

- Use of an *interface class* allows `callback` to be treated as *mixin* functionality
  - note that `callback_client`'s destructor is *protected* and *non-virtual*, clearly stating that it is not intended as primary design type

```
class rule_base
{
public:
    void for_each(
        callback_client& c) const;
    ...
};
```

```
class callback_client
{
public:
    virtual void
        callback(const rule& obj) = 0;
protected:
    ~callback_client();
};
```

```
class rules_view :
    public view,
    private callback_client
{
public:
    void update(const rule_base& rb)
    { rb.for_each(*this); }
    ...
private:
    virtual void callback(
        const rule& r)
    { Code to add rule to view ... }
    ...
};
```

19

## The Problem of Selective Traversal

- In many cases the callback function should only be called for `rules` that execute in response to a particular `trigger` event
  - Making filtering options available via overloading means the choice of filtering must be made at compile time

```
class rule_base
{
public:
    void for_each(callback_client& c) const;
    void for_each(callback_client& c, const trigger& e) const;
    ...
};
```

20

## Adding a Run Time Filter

```
class callback_filter
{public:
    virtual bool required(const rule& obj)
        const=0;
protected:
    ~callback_filter();
};
```

```
class rule_base
{public:
    void for_each(
        callback_client& client,
        const callback_filter& filter) const
    {
        for (const_it it = rules.begin();
            it != rules.end();
            ++it)
            if (filter.required(**it))
                client.callback(**it);
        ...
    };
```

- The most versatile approach involves `for_each()` taking an interface class that defines a (virtual) filter function as a second parameter
  - The filter function returns `true` if the `rule` is to be included

21

## Tradeoffs of Run Time Filter Approach

- Introducing run-time polymorphism affords more versatility than the compile-time polymorphism of overloading
- But, `rule_base`'s implementation can't be taken advantage of
  - E.g. maintaining indices to improve traversal performance is much more complicated

22

### **Design Decision Re-evaluated**

- With the benefit of hindsight, using a Batched Copying approach was probably simpler
  - It isn't clear that the framework afforded by Enumeration by Callback (in conjunction with Callback Using Interface Class) lead to cleaner code

23

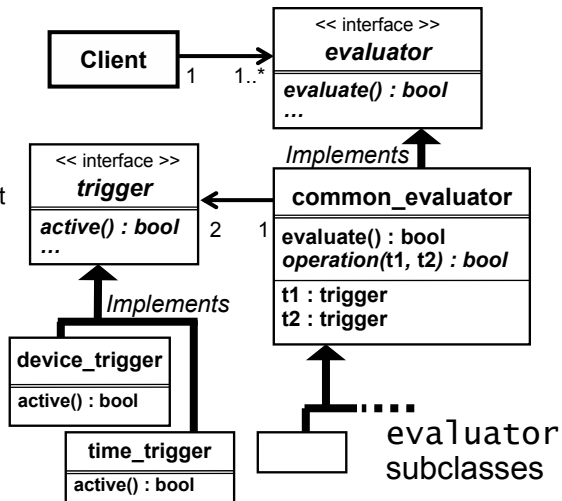
### **Mixing Generic with Object Oriented Programming**

- Design of (part of) a logical expression evaluation engine
- The tradeoffs in mixing dynamic polymorphism with generic programming in design

24

## An Evaluation Engine for a Security System

- trigger objects represent:
  - Physical devices such as motion detectors
  - Time intervals that are entered daily, weekly etc.
- evaluator evaluates whether or not a combination of triggers requires action to be taken



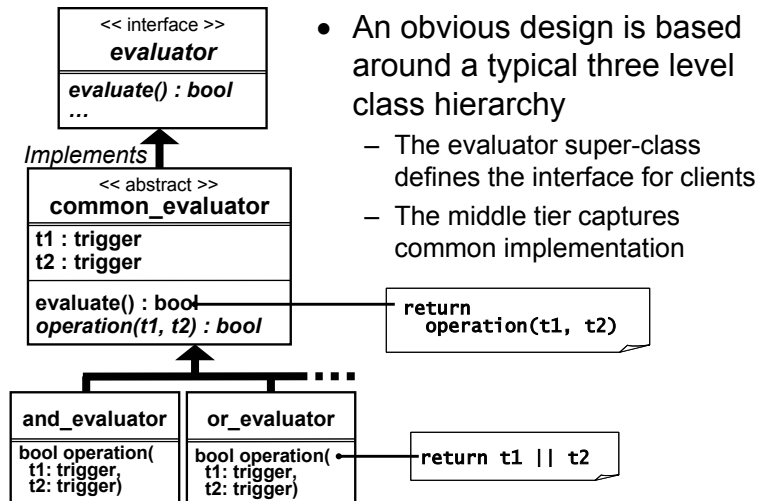
25

## The C++ Three Level Hierarchy

- The three level inheritance hierarchy is a common design idiom in C++
  - The base class implements only the common interface (pure virtual functions only)
  - The next level of derivation captures common implementation detail in an abstract class
  - The most derived level contains the concrete classes

26

## The evaluator Hierarchy



27

## The Principle C++ *Interface* Classes

```

class trigger
{
public:
    virtual ~trigger();
    virtual bool active() const = 0;
};
    
```

```

class evaluator
{
public:
    evaluator() {}
    virtual ~evaluator();
    virtual bool evaluate() const = 0;
private:
    evaluator(const evaluator&);
    evaluator& operator=(const evaluator&);
};
    
```

- trigger & evaluator are *interface* classes in C++
  - trigger's active() member function returns true if the device the trigger object represents has been activated
  - evaluator's evaluate() evaluates an expression made up of two trigger objects

28

## The Problem of Repetitive Code

```
class common_evaluator : public evaluator
{protected:
    common_evaluator(
        shared_ptr<trigger> trigger_1, shared_ptr<trigger> trigger_2);
private:
    virtual bool evaluate() const { return operation(*t1, *t2); }
    virtual bool operation(const trigger& t2, const trigger& t1) const = 0;
    trigger::shared_ptr t2, t1;
    ...
};
```

```
class or_evaluator : public common_evaluator
{public:
    or_evaluator(
        boost::shared_ptr<trigger> trigger_1,
        boost::shared_ptr<trigger> trigger_2)
        : common_evaluator(trigger_1, trigger_2)
    {}
private:
    virtual bool operation(
        const trigger& trigger_1,
        const trigger& trigger_2) const
    { return trigger_1.active() ||
        trigger_2.active(); }
    ...
};
```

- The repetition of code in the concrete classes introduces a risk of error
  - Note the erroneous appearance of `trigger_1` which may or may not have been duplicated in `and_evaluator`

29

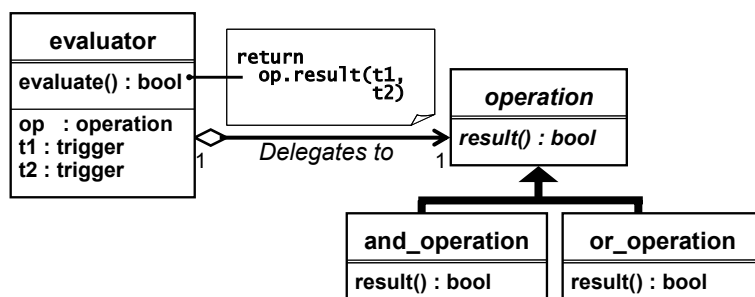
## Alternative Factorings of Variability

- An alternative factoring taking better account of the commonality/variability in the `evaluator` hierarchy is needed
  - The variability is purely behavioural (i.e. the relational operation)
- Options are available for configuring the variability at both statically and dynamically
  - Naturally each has different tradeoffs associated with it

30

## Solution Using Run-Time Delegation

- The variability can be factored into a separate *interface* class
  - Evaluator is then configured at when it is instantiated



31

## Delegation Solution – Implementation

```

class operation
{public:
    virtual bool result(
        const trigger& t1, const trigger& t2)
        const = 0;
    ...
};
    
```

```

class evaluator
{public:
    evaluator(
        boost::shared_ptr<trigger> trig_2,
        boost::shared_ptr<trigger> trig_1,
        boost::shared_ptr<operation> eval_op);
private:
    virtual bool evaluate() const
    { return op->result(*t1, *t2); }

    boost::shared_ptr<trigger> t2,
    boost::shared_ptr<trigger> t1,
    boost::shared_ptr<operation> op;
};
    
```

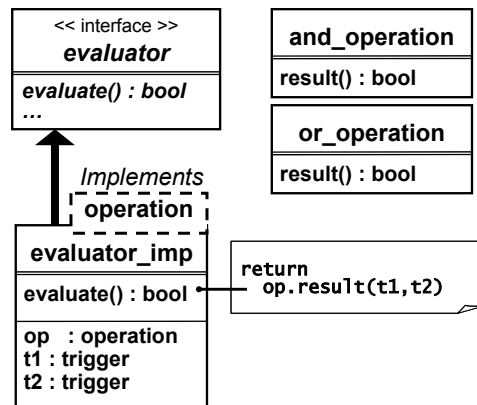
- Less repetition of code leads to fewer opportunities for error
- But there are still many objects whose lifecycles must be managed
  - The client code has to deal with operations as well as triggers

32



## Solution Using *Static* Parameterisation

- The variability can be captured using a parameterised class
  - A C++ *class template* being the natural implementation



33

## Implementation by C++ Class Template

```

template <typename function>
class evaluator_implementor : public evaluator
{public:
    evaluator_implementor(
        shared_ptr<trigger> trig1, shared_ptr<trigger> trig2)
        : t1(trig1), t2(trig2) {}
private:
    virtual bool evaluate() const
    { return function()(t1->active(), t2->active()); }
    shared_ptr<trigger> t1, t2;
};

typedef evaluator_implementor<std::logical_and<bool> >
    and_evaluator;
...
  
```

- Specialisations for *and/or* (etc.) can be created using function objects available in the standard library

34

## Observations

- An approach using delegation leads to simpler and less error prone code
  - The approach using inheritance would produce lots of similar code
- Combining generics with dynamic polymorphism lead to the simplest implementation
  - No duplication of almost the same code
  - Easy to extend (e.g. for an XOR operation)

35

## Complexity & Design

- Ignoring complexity is costly
  - How to deal with the complexity is an issue to be addressed during design
- C++ has features designed to help absorb complexity
  - E.g. templates and exceptions
  - But understanding these (advanced) features places an added burden on the skills of the programmer

36

## **Observations on Design**

- Design/programming paradigms
- Multi-Paradigm design
- An alternative perspective on combining design/programming paradigms

37

## **Inheritance & Run Time Polymorphism**

- It is a popular myth that inheritance is a defining feature of an object oriented technology
  - Several programming languages have helped propagate this myth by using inheritance as their run-time polymorphism mechanism
- Inheritance and run-time polymorphism are individual and separate concepts

38

## Object Design Paradigms

- Object based
  - Supports abstraction and encapsulation
  - E.g. value based programming using the *Whole Value* idiom
- Object oriented
  - Defined by the presence of run-time polymorphism as a first class feature

39

## Other Design Paradigms

- Procedural
  - Concerned with flow of control
- Generic
  - Separates types, data structures, operations and flow of control
  - Not useful in its own right, but used in combination with other paradigms

40

## Multi-Paradigm Design

```
class shape
{public:
    virtual void rotate(const radians& angle) = 0;
    ...
};
class line : public shape { ... };
...
```

```
void rotate_all(
    std::vector<shape*>& shapes,
    const radians& angle)
{
    typedef std::vector<shape*>::iterator
        iterator;
    for (iterator pos = shapes.begin();
        pos != shapes.end();
        ++pos) {
        shape* current = *pos;
        current->rotate(angle);
    }
}
```

Generic  
programming

Procedural  
programming

Object oriented  
programming

41

## Tools in the Mental Toolbox

```
class shape
{public:
    virtual void rotate(const radians& angle) = 0;
    ...
};
class line : public shape { ... };
...
```

```
void rotate_all(
    std::vector<shape*>& shapes,
    const radians& angle)
{
    typedef std::vector<shape*>::iterator
        iterator;
    for (iterator pos = shapes.begin();
        pos != shapes.end();
        ++pos) {
        shape* current = *pos;
        current->rotate(angle);
    }
}
```

Static  
parameterisation

Flow of control

Run-time  
polymorphism

42

## Designing With Objects

- One of the great benefits of C++ is its support for a variety of design paradigms
  - But in practice, distinguishing between paradigms constrains the thought process
  - It is better to view elements of paradigms as tools in the intellectual toolbox of the design process as a whole

43

## The End

- I hope you found this talk interesting
- Thank you for your attention!

44